


UNDERSTAND IT IN 30 SECONDS

Race evidence in. Reviewed stories out.

A race-reporting workspace that brings sources, cited drafts and editorial decisions together.

[Explore the product](#)[Jump to the evidence](#)

Meet the workflow first. Inspect every claim next.

WHO CAN USE IT

For the people behind every race update.

Designed for teams that need to explain what changed, show the source and keep an editor in control.

01 Race organizers

Prepare event updates from timing facts and course context.

Clearer event communications

02 Sports editors

Inspect citations and revise the story before approving a version.

A reviewable editorial draft

03 Event media teams

Start with a structured story and a short social caption.

A shared starting point

04 Running communities

Explain position, pace, cutoff and record-projection signals.

Race moments with context

The product promise


One place to move from a race signal to a sourced draft, a documented check and an editorial decision.

Intended audiences and benefits; no customer adoption, time savings or revenue claim is implied.

HOW IT WORKS

A visible path from signal to editorial decision.

The model prepares a candidate. Validation and human review determine what happens next.




Illustrative example, not a saved model output or real race result.

Editor control: source inspection, protected revisions and history. Approval does not publish externally.

THE WEEK 5 LEARNING STORY

Teach the response behavior. Supply the race facts.

Prompting defines the task; retrieval supplies current evidence; fine-tuning adapts the response behavior.



A dataset with a separate test set



150 fictional race groups; no group crosses splits. Four signal types, five evidence conditions.

Demonstrated

Qwen3-4B-Instruct-2507; rank-16 LoRA.
NF4 training: 2 epochs, 100 steps.
33,030,144 trainable parameters.
Adapter merged and evaluated locally.

Future experiments / pending proof


Fireworks RFT and teacher distillation: not run.
Native Ollama API demo: not demonstrated.
Human editorial benefit: not measured.
Full agenda mapping is linked on page 6.

All training examples are programmatic synthetic references; none are human-approved.

WHAT THE DATA ACTUALLY SHOWS

A measured gain, with the limits beside it.

Matched local comparison: same base lineage, prompt, Q4_K_M runtime and schema constraints.



+27.5 points

Paired 95% interval: +20 to +35 points

What passed

Both models: 120/120 valid schemas.
The score combines eight checks for structure, citations, claims, numeric support, decision and selected wording.

Evidence that keeps the claim in bounds

11 fine-tuned failures: unnecessary holds.
Five new probes: base 3/5; adapted 3/5.
Rules: 120/120 on the engineered test.

13 app software checks passed; separate quality suite failed one case. Human review and the paired explicit-schema GPU control remain pending.

Synthetic benchmark only; not editor preference. Rules share the task design. Bootstrap: 30 groups, 1,000 resamples.

GO DEEPER AT YOUR OWN PACE

Start with the demo. Follow the evidence.


Every underlined resource and colored button is clickable. QR codes also work from a printed copy.



Explore UltraMedia

Product and editorial workflow

Open online



Open the evidence hub

Week 5 results and recorded examples

Open online

01 Product walkthrough

Watch the recorded workflow.

03 Data and provenance

Split design, examples and hashes.

05 Paired model comparison

Scores, raw outputs and failures.

07 Week 5 agenda map

Executed topics versus future work.

02 Complete project report

Requirements, tradeoffs and decisions.

04 Training receipts

Configuration, loss and verification.

06 Case-level CSV

Inspect all paired test rows.

08 End-to-end visual

See the full technical illustration.

Public demo includes recorded examples. This brochure does not claim live cloud model hosting or production approval.